



Stable Fluids

DAVID PERRONE
JAKE AFONSO

Presentation overview

- ▶ Algorithm
- ▶ Parallelization
- ▶ Rendering
- ▶ Demo

Background

- ▶ Simulates a square of a fluid, with a defined viscosity and diffusion
 - ▶ Viscosity is the thickness of the fluid
 - ▶ Diffusion is how fast density spreads through the fluid
- ▶ A “dye” is added at some point of the square with a given velocity
- ▶ Dye moves throughout the fluid based on the fluid dynamics defined in the Navier-Stokes equations
- ▶ Stable in the fact that there are no bounds on the time steps and the algorithms never “blow up”

Diffusion


- ↖ Represents the property of a dye spreading out throughout a fluid.
- ↖ Happens for the dye itself and less obviously for the velocities of the fluid
- ↖ Can happen differently based on the viscosity/diffusion value of the liquid
 - ↖ For example, putting dye in water will spread much faster than in molasses
 - ↖ For our simulation, velocity diffusion is based on the viscosity value and density diffusion is based on the diffusion value
- ↖ Each particle exchanges density and velocity with its 4 direct neighbors



Start




End



Advection

- ↖ Represents the movement of the dye based on the velocity field within the fluid
- ↖ If a liquid is moving, it will carry the dye in that direction as well
- ↖ Similarly to diffusion, this happens for the dye and the velocities themselves



Advection carries the dye through the velocity field

Projection

- ↖ This upholds the incompressible property of the fluid
- ↖ The amount of fluid coming into the square must equal the amount of fluid leaving the square
- ↖ The other 2 functions mess with the equilibrium of the fluid, and project cleans it up after



Subroutines

Linear Solve

All 3 of these functions are based on approximating a solution to a linear system

Implements the Gauss-Seidel method → iterative solver

Set bnd

These functions all only deal with the inner part of the square

Set boundary sets the edges to be the opposite of the neighboring inner particle

Gives the effect of bouncing off the walls

Parallelization

- ▶ Used cuda to implement a parallel version of the simulation
- ▶ Defined num_threads/num_blks to have each thread be assigned a single particle
- ▶ Diffuse, advect, and project functions into __global__ functions to launch as kernels
- ▶ Lin_solve and set_bnd into __device__ functions to be called from within the kernels
- ▶ Performed slower than serial for $n < 6400$ particles, but then started to outperform for $n > 6400$.
 - ▶ Overall close to constant time for increasing size, while serial was exponential

Serial:

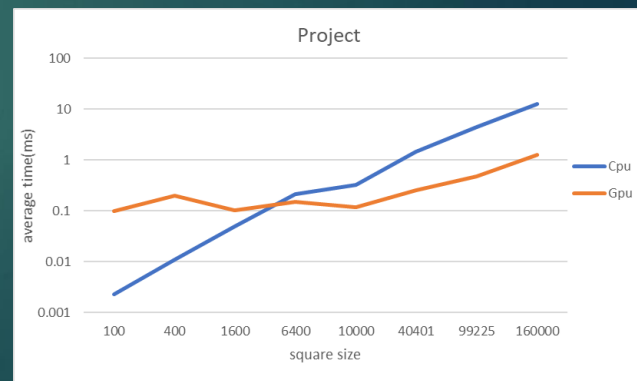



```
void lin_solve(int b, float* x, float* x0, float a, float c, int iter, int N) {  
  
    float cRecip = 1.0 / c;  
  
    for (int k = 0; k < iter; k++) {  
        for (int j = 1; j < N - 1; j++) {  
            for (int i = 1; i < N - 1; i++) {  
                x[IX(i, j)] =  
                    (x0[IX(i, j)]  
                     + a * (x[IX(i + 1, j)]  
                          + x[IX(i - 1, j)]  
                          + x[IX(i, j + 1)]  
                          + x[IX(i, j - 1)]  
                          )) * cRecip;  
            }  
        }  
        set_bnd(b, x, N);  
    }  
}
```

Cuda:

```
__device__ void lin_solve_gpu(int b, float* x, float* x0, float a, float c, int iter, int N, int tid)  
{  
    int localID = threadIdx.x;  
    int i = tid % N;  
    int j = tid / N;  
  
    __shared__ float local_x[NUM_THREADS];  
  
    if (i < 1 || i > N - 2) return;  
    if (j < 1 || j > N - 2) return;  
  
    float cRecip = 1.0 / c;  
    for (int k = 0; k < iter; k++) {  
        local_x[localID] =  
            (x0[IX(i, j)]  
             + a * (x[IX(i + 1, j)]  
                  + x[IX(i - 1, j)]  
                  + x[IX(i, j + 1)]  
                  + x[IX(i, j - 1)]  
                  )) * cRecip;  
        __syncthreads();  
        x[IX(i, j)] = local_x[localID];  
        __syncthreads();  
        set_bnd_gpu(b, x, N, tid);  
    }  
}
```

```
void set_bnd(int b, float* x, int N) {  
    // set velocity at boundary to be opposit of the velocity for the tile ne  
    for (int i = 1; i < N - 1; i++) {  
        x[IX(i, 0)] = b == 2 ? -x[IX(i, 1)] : x[IX(i, 1)];  
        x[IX(i, N - 1)] = b == 2 ? -x[IX(i, N - 2)] : x[IX(i, N - 2)];  
        x[IX(0, i)] = b == 1 ? -x[IX(1, i)] : x[IX(1, i)];  
        x[IX(N - 1, i)] = b == 1 ? -x[IX(N - 2, i)] : x[IX(N - 2, i)];  
    }  
  
    x[IX(0, 0)] = 0.5 * (x[IX(1, 0)] + x[IX(0, 1)]);  
    x[IX(0, N - 1)] = 0.5 * (x[IX(1, N - 1)] + x[IX(0, N - 2)]);  
    x[IX(N - 1, 0)] = 0.5 * (x[IX(N - 2, 0)] + x[IX(N - 1, 1)]);  
    x[IX(N - 1, N - 1)] = 0.5 * (x[IX(N - 2, N - 1)] + x[IX(N - 1, N - 2)]);  
}
```

```
__device__ void set_corner_gpu(float* x, int N)  
{  
    x[IX(0, 0)] = 0.5 * (x[IX(1, 0)] + x[IX(0, 1)]);  
    x[IX(0, N - 1)] = 0.5 * (x[IX(1, N - 1)] + x[IX(0, N - 2)]);  
    x[IX(N - 1, 0)] = 0.5 * (x[IX(N - 2, 0)] + x[IX(N - 1, 1)]);  
    x[IX(N - 1, N - 1)] = 0.5 * (x[IX(N - 2, N - 1)] + x[IX(N - 1, N - 2)]);  
}  
  
__device__ void set_bnd_gpu(int b, float* x, int N, int tid)  
{  
    int i = tid % N;  
    int j = tid / N;  
    if (i >= 1 && i < N - 1 && j >= 1 && j < N - 1)  
    {  
        x[IX(i, 0)] = b == 2 ? -x[IX(i, 1)] : x[IX(i, 1)];  
        x[IX(i, N - 1)] = b == 2 ? -x[IX(i, N - 2)] : x[IX(i, N - 2)];  
        x[IX(0, j)] = b == 1 ? -x[IX(1, j)] : x[IX(1, j)];  
        x[IX(N - 1, j)] = b == 1 ? -x[IX(N - 2, j)] : x[IX(N - 2, j)];  
    }  
    __syncthreads();  
    int n2 = N / 2;  
    if (i == n2 && j == n2) {  
        set_corner_gpu(x, N);  
        return;  
    }  
}
```




Rendering

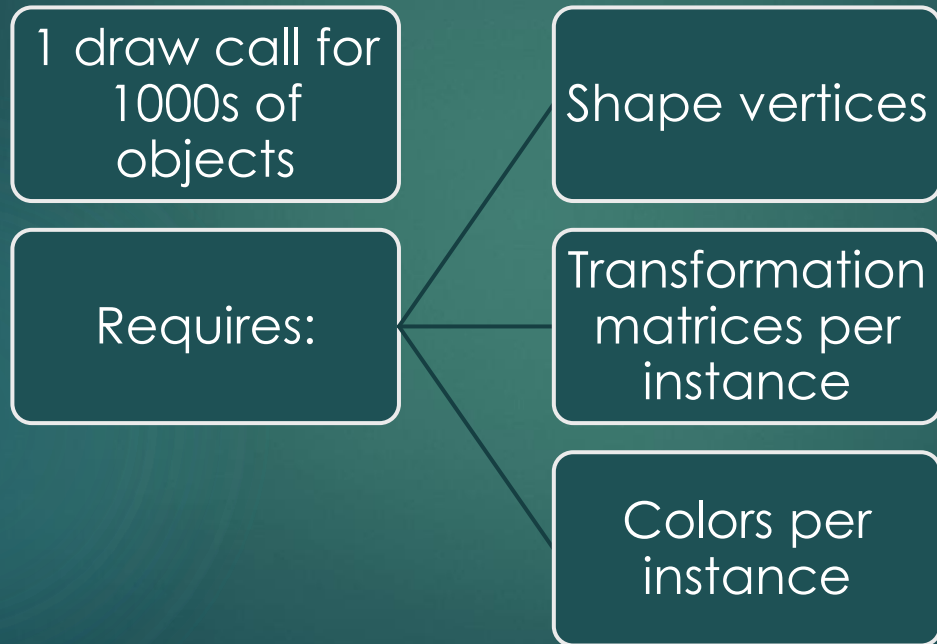
- ↳ Vertices
- ↳ Instanced Rendering
- ↳ Cuda OpenGL Interop
- ↳ OpenGL Memory layout
- ↳ Matrix Multiplication

Rendering in OpenGL

- ▶ The fluid square is made up of several smaller squares
- ▶ Each square is rendered as two triangles
- ▶ Multiply each vertex of the square with a transformation matrix to move/scale
- ▶ The color within the square is represented by 4 floats (RGBA)



Instanced Rendering




Memory

Cuda performs
the square
calculations

- Stores the location for each square
- Stores the density (color) for each square

OpenGL
renders the
updated data
each frame

Data Transfer



Rearranging the memory

Transformed
vertices

Color




vertices



Colors



Transformation
Matrices




Cuda - OpenGL Communication


1. Create CudaGraphicsResource
2. Create OpenGL Vertex Array Object
3. Create OpenGL vertex buffer - fill with 4 vertices of square
4. Create OpenGL Matrix Buffer
5. Register CudaGraphicsResource with Matrix Buffer
6. Create translation matrices, multiply with one scale matrix - store into Buffer
7. UnRegister Cuda
8. Repeat with OpenGL Color Buffer
9. NOTE: OpenGL and Cuda cannot use the same buffer simultaneously

Matrix Multiplication


Transformation
Matrices



Scale Matrix

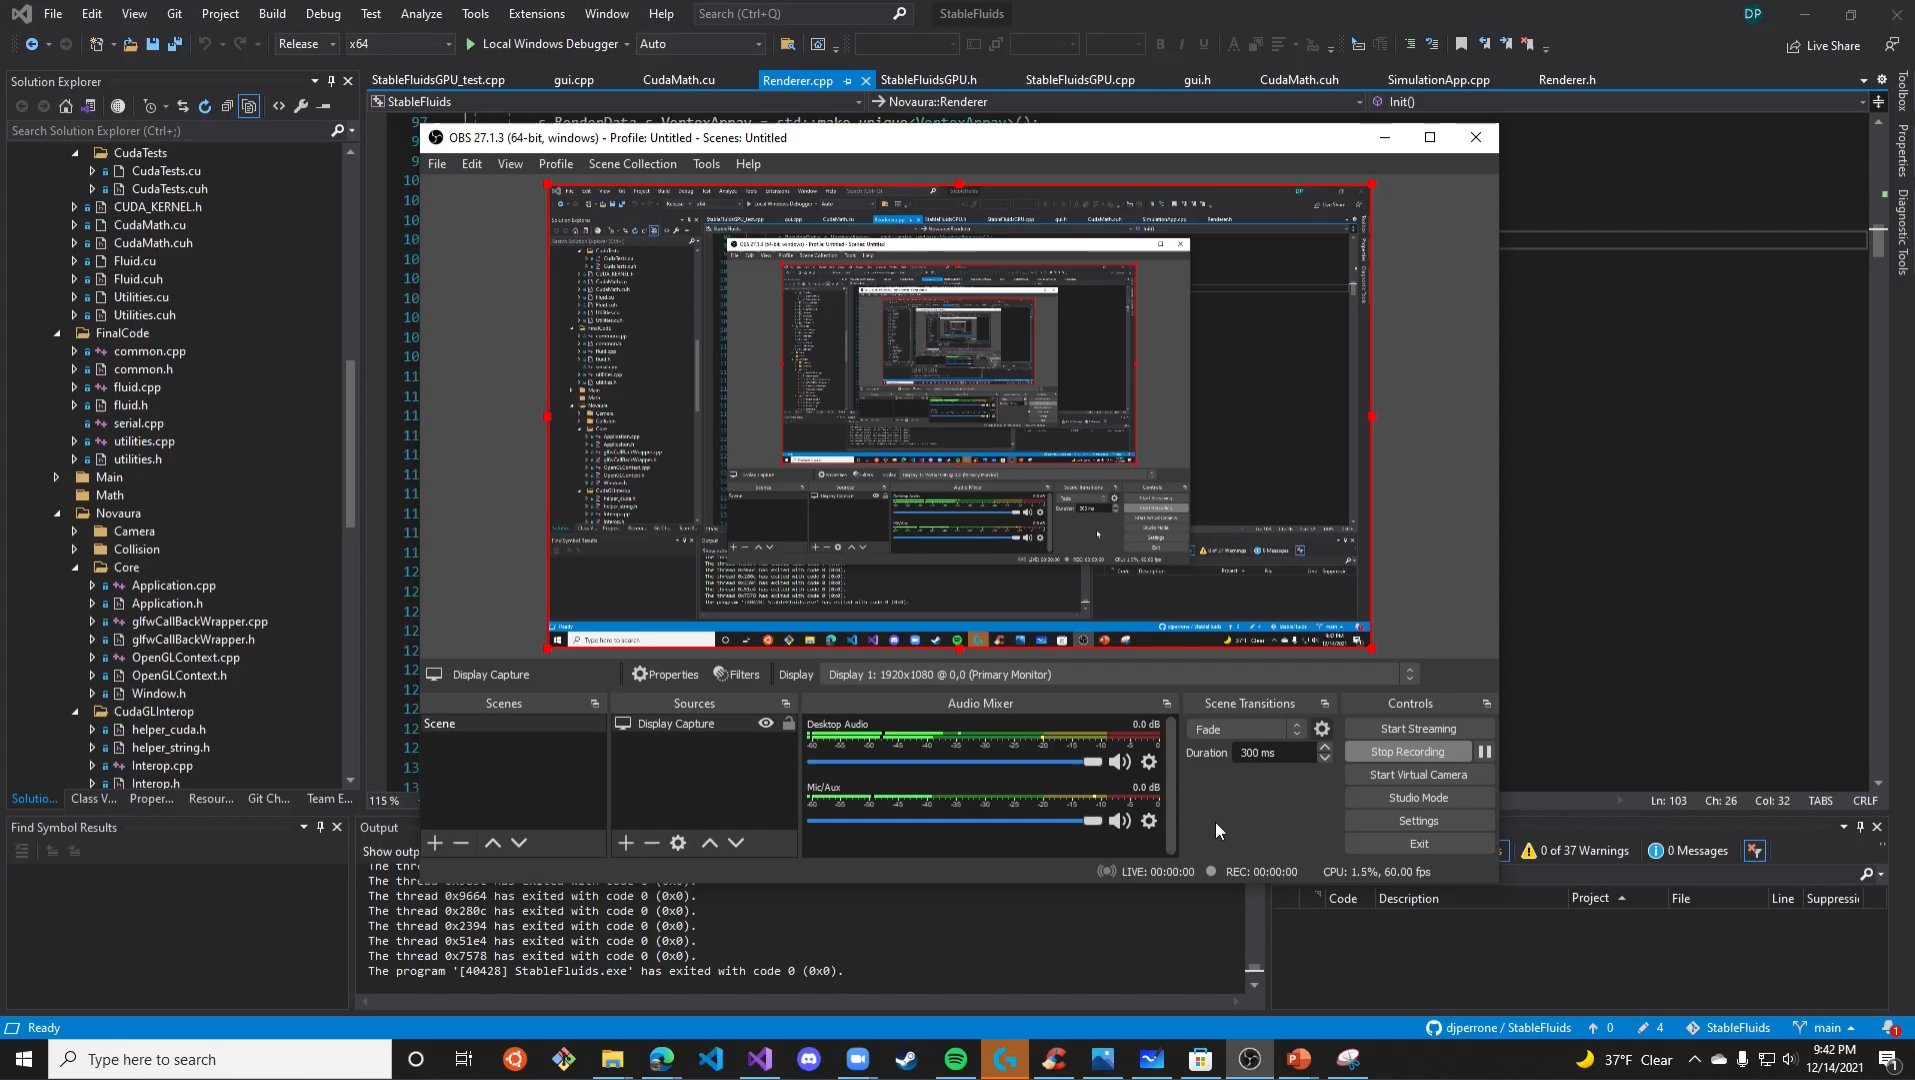


OpenGL Matrices



Demo (Mute Audio)





Questions?



Github

↴ [djperrone/StableFluids \(github.com\)](https://github.com/djperrone/StableFluids)